



# k Nearest Neighbors algorithm

# Outline

## ● KNN

- Background
- Definition
- Distance measures
- How to choose  $k$
- Digit recognition
- K-nn properties

## ● K-D TREE CONSTRUCTION & QUERY

## ● HOMEWORK

---

## Origins of K-NN

- Nearest Neighbors have been used in statistical estimation and pattern recognition in the beginning of 1970's (non-parametric techniques).
  - The method prevailed in several disciplines and still it is one of the **top 10 Data Mining algorithm**.
-



You are the average of the five people you spend most time with.

—Jim Rohn


- It's how people judge by observing our peers.



- We tend to move with people of similar attributes so does data

# Definition

- K-Nearest Neighbor is considered a lazy learning algorithm that classifies data sets based on their similarity with neighbors.
- “K” stands for number of data set items that are considered for the classification.



Note: K -Nearest Neighbors is called a non-parametric method  
Unlike other supervised learning algorithms, K -Nearest Neighbors doesn't learn an explicit mapping  $f$  from the training data  
It simply uses the training data at the test time to make predictions

## Technically

- Given training data  $\{(x_1, y_1), \dots, (x_N, y_N)\}$  and a test point  $x$ 
  - $N$  pairs ;  $x_i$  is a **vector** consisting of  $D$  features ,  $y_i$  - label
- Goal: predict the output  $y$  for an unseen test example  $x$
- Prediction rule: look at the  $K$  most similar training examples

# Technically

- Forms of the output:

for classification: a discrete variable  $y_i \in \{1, \dots, C\}$   
assign the majority class label (**majority voting**)

for regression: a continuous (real-valued) variable  $y_i \in R$   
assign the **average response**

- The algorithm requires:

**Parameter K**: number of nearest neighbors to look for

**Distance function**: To compute the similarities between examples

---



## K-NN: Some distance measures

- Euclidean distance is commonly used

In the Euclidean plane, if  $\mathbf{p} = (p_1, p_2)$  and  $\mathbf{q} = (q_1, q_2)$  then the distance is given by

$$d(\mathbf{p}, \mathbf{q}) = \sqrt{(q_1 - p_1)^2 + (q_2 - p_2)^2}.$$

In three-dimensional Euclidean space, the distance is

$$d(p, q) = \sqrt{(p_1 - q_1)^2 + (p_2 - q_2)^2 + (p_3 - q_3)^2}.$$


In general, for an  $n$ -dimensional space, the distance is

$$d(p, q) = \sqrt{(p_1 - q_1)^2 + (p_2 - q_2)^2 + \cdots + (p_i - q_i)^2 + \cdots + (p_n - q_n)^2}.$$

# distance measures

- **Manhattan Distance**(or City Block distance)

$$d(\mathbf{p}, \mathbf{q}) = |p_1 - q_1| + |p_2 - q_2|$$



In general, for an  $n$ -dimensional space, the Manhattan distance is

$$d(\mathbf{p}, \mathbf{q}) = \|\mathbf{p} - \mathbf{q}\| = \sum_{i=1}^n |p_i - q_i|$$

# distance measures

- **Minkowski Distance**(闵可夫斯基距离)

The Minkowski distance of order  $p$  between two points

$$p = (p_1, p_2, \dots, p_n) \text{ and } q = (q_1, q_2, \dots, q_n) \in \mathbb{R}^n$$

is defined as:

$$\left( \sum_{i=1}^n |p_i - q_i|^m \right)^{1/m}$$

Minkowski distance is typically used with  $m$  being 1 or 2, which correspond to the Manhattan distance and the Euclidean distance, respectively.

In the limiting case of  $m$  reaching infinity, we obtain the Chebyshev distance.

## distance measures

- **Cosine similarity** is a measure of similarity between two non-zero vectors of **an inner product space** that measures the **cosine of the angle** between them.

the cosine similarity,  $\cos(\theta)$  is represented using:

$$\text{similarity} = \cos(\theta) = \frac{A \cdot B}{\|A\| \|B\|} = \frac{\sum_{i=1}^n A_i \times B_i}{\sqrt{\sum_{i=1}^n (A_i)^2} \times \sqrt{\sum_{i=1}^n (B_i)^2}}$$

## distance measures

- **Hamming distance** (汉明距离) counts the number of features where the two examples disagree.

The Hamming distance between:

• **1011101** and **1001001** is 2.

• **2173896** and **2233796** is 3.

In other words, it measures the minimum number of *substitutions* required to change one string into the other .

# K -NN: Feature Normalization

Note: Features should be on the same scale

Example: if one feature has its values in millimeters and another has in centimeters, we would need to normalize

One way is:

Replace  $p_i$  by  $z_{im} = \frac{(x_{im} - \bar{x}_m)}{\sigma_m}$  (make them zero mean, unit variance)

$\bar{x}_m = \frac{1}{N} \sum_{i=1}^N x_{im}$  : mean of  $m^{th}$  feature


$\sigma_m^2 = \frac{1}{N} \sum_{i=1}^N (x_{im} - \bar{x}_m)^2$  : variance of  $m$  feature

# K-NN: Feature Weighting

- Scale each

- Can use ou  
important

- Can learn t



ation

res are more

(not covered)

How about the weight of the sample?

## K-NN: How to Choose k?

- In theory, if **infinite** number of samples available, **the larger** is k, **the better** is classification.
- Impossible in practice since # samples is **finite**


□  $k=1$     the nearest one

□  $k=N$     the majority




# K-NN: How to Choose k?

- $k = 1$  is often used for efficiency, but sensitive to “noise”




every example in the blue shaded area will be misclassified as the blue class



every example in the blue shaded area will be classified correctly as the red class

# 1NN Visualization


- **Voronoi Diagram** (维诺图) is useful for visualization



A Voronoi diagram divides the space into such cells.

Each cell contains one sample, and every location within the cell is closer to that sample than to any other sample.

# Choice of K - Neighborhood Size



Small K

Creates many small regions for each class  
sensitive to “noise”

(May lead to non-smooth) decision boundaries and overfitting

Large K


creates fewer larger regions ,  
usually produces smoother decision boundaries ,  
can reduce the impact of class label noise.

(caution: too smooth decision boundary can underfitting)

# Holdout

- Divide labeled data into two sets
  - A training set (80%) and a validation set (20%)
- Predict the class labels for validation set by using the examples in training set
- This is an estimate of the generalization error
- Choose the number of neighbors  $k$  that maximizes the classification accuracy
- Once we have chosen the model, and report performance on a test set.

# K-fold cross-validation



Final Validation Accuracy =  $\text{mean}(\text{val\_acc1} + \text{val\_acc2} + \dots + \text{val\_acc10})$

# Leave-One-Out Method

- For  $k = 1, 2, \dots, K$

$$err(k) = 0$$

For  $i = 1, 2, \dots, n$

- \* Predict the class label  $\hat{y}_i$  for  $\mathbf{x}_i$   
using the remaining data points
  - \*  $err(k) = err(k) + 1$  if  $\hat{y}_i \neq y_i$
- Output  $k^* = \arg \min_{1 \leq k \leq K} err(k)$

# K-NN: How to Choose?

- What distance measure to use?

Often **Euclidean** distance is used

Locally adaptive metrics

More complicated with non-numeric data, or when different dimensions have different scales

- Choice of  $k$ ?

odd number

1-NN often performs well in practice

Interesting theoretical properties if  $k < \sqrt{n}$ ,  $n$  is # of examples

Can choose  $k$  through **cross-validation and so on**


## Example: Digit Recognition



**Kaggle** is a platform for predictive modelling and analytics competitions to produce the best models for predicting and describing the data.




# Example: Digit Recognition

kaggle


17 active competitions

All Categories


Search competitions



**Passenger Screening Algorithm Challenge**  
Improve the accuracy of the Department of Homeland Security's threat recognition algorithms  
**Featured** · 3 months to go



**Zillow Prize: Zillow's Home Value Prediction (Zestimate)**  
Can you improve the algorithm that changed the world of real estate?  
**Featured** · 4 months to go



**Carvana Image Masking Challenge**  
Automatically identify the boundaries of the car in an image  
**Featured** · 22 days to go

**\$1,500,000**  
241 teams

**\$1,200,000**  
2,653 teams

**\$25,000**  
549 teams

## Example: Digit Recognition

kaggle™

9 6 6 5 4 0 7 4 0 1  
3 1 3 4 7 2 7 1 2 1  
1 7 4 2 3 5 1 2 4 4

### Digit Recognizer

Learn computer vision fundamentals  
with the famous MNIST data  
1,573 teams · 2 years to go

Type	Classifier	Preprocessing	Error rate (%)
Convolutional neural network	Committee of 5 CNNs	Expansion of the training data	0.21
Convolutional neural network	Committee of 35 CNNs	Width normalizations	0.23
Deep neural network	6-layer	None	0.35
K-Nearest Neighbors	K-NN with non-linear deformation	Shiftable edges	0.52
Support vector machine	Virtual SVM	Deskewing	0.56
Non-linear classifier	PCA + quadratic classifier	None	3.3
Linear classifier	Pairwise linear classifier	Deskewing	7.6

Nearest neighbour is still fairly competitive!

# Example: Digit Recognition



Number Plate  
Recognition



House no. Recognition


中国工商银行 INDUSTRIAL AND COMMERCIAL BANK OF CHINA		储蓄 存款 凭条 SLIP FOR SAVINGS DEPOSIT	
科目: (贷)		2001年 6月20日	
银行 填写	*007100692660	2001-06-20	01-06-20
包月券 1101	1	02-06-20	0 000000000000
交易代码:		047*01	1
存款: 整存 <input type="checkbox"/> 零存 <input type="checkbox"/> 活期 <input type="checkbox"/> 定活 <input type="checkbox"/> 通知 <input type="checkbox"/> 整存 <input type="checkbox"/> 存本 <input type="checkbox"/> 教育 <input type="checkbox"/> 零存 <input type="checkbox"/> 零通 <input type="checkbox"/> 零通 <input type="checkbox"/> 其他 <input type="checkbox"/>		3,000.00 00755	
储户 填写	户名 包月芳	帐号	存款 1 币种 人民币
身份证件名称	号码 511021145170110613731	金额	叁仟零元整
代理身份证件名称	号码 511021145170110613731	利息	300000
新开户填写: 地址		印密 <input type="checkbox"/> 通存通兑 <input type="checkbox"/> 约定转存 <input type="checkbox"/>	事后监督 复核(授权) 经办

Recognition of legal  
amounts on bank cheques

FROM: John Smith 12 Main Street Mytown, PE C1A 2J2 Canada	TO: Jane Brown 10 Some Street Janestown, NY 12345 USA
---	---

Postal mail sorting

# Data Set



- ▶ testDigits
- ▼ trainingDigits
  - 0\_0.txt
  - 0\_1.txt
  - 0\_2.txt
  - 0\_3.txt
  - 0\_4.txt
  - 0\_5.txt
  - 0\_6.txt
  - 0\_7.txt
  - 0\_8.txt
  - 0\_9.txt

...

[illegible]

32x32 pixel images:  $d = 1024$   
1,934 training samples  
946 test samples

# Example: Digit Recognition


## 机器学习 K-NN手写识别 demo

### 简介

这个演示使用K-NN算法，为了简便，只能识别数字0-9。

数据集来自[Machine Learning in Action](#) 第二章。书中把来自[UCI数据库](#)的手写数据集简化成32像素x32像素的黑白图像，并且以01矩阵的方式存储在txt文件中。大约有训练样本2000个，测试样本900个。我使用python脚本将原始的txt文件转化成Json格式，以便减少文件下载的请求次数。

### 测试



清空

手写识别

集合测试

☒ KNN

☐ CNN

请输入K：

20


### 结果

<http://211.87.235.83/other/HandwritingRecognition/>


# Curse of Dimensionality

- K-NN breaks down in high-dimensional space, “**Neighborhood**” becomes very large and easily **misleading**.
- The curse of dimensionality refers to **various** phenomena that arise **in high-dimensional spaces** that do not occur in low-dimensional settings of everyday experience.
  - Storage complexity
  - Computational complexity
  - Sampling
  - Combinatorics
  - Nearest neighbor search
  - Distance functions
  - Nonparametric estimation
  - ...


# sampling



1维: 5





2维:  $5 \times 5 = 25$



10维:  $5^{10} = 976\,5625$

# Nearest neighbor search

- Assume 5000 points uniformly distributed in the unit hypercube and we want to apply 5-nn. (Suppose our query point is at the origin)
  - In 1-dimension, we must go a distance of  $5/5000 = 0.001$  on the average to capture 5 nearest neighbors
  - In 2 dimensions, we must go  $\sqrt{0.001}$  to get a square that contains 0.001 of the volume.
  - In  $d$  dimensions, we must go  $(0.001)^{1/d}$





# 维数灾难的几个表现

## ○ 噪声影响

- 特征空间：101维
- 正负样本在第一维的距离：1
- 样本在其余维的噪声：10%
- “噪声距离”： $\sqrt{100 \times 0.1^2} = 1$
- □即使噪声只有10%，高维空间的“噪声距离”足以掩盖正负样本的本质区别

<http://www.cnblogs.com/zhangchaoyang/articles/2801525.html>

<http://blog.csdn.net/zc02051126/article/details/49618633>

<https://zh.wikipedia.org/wiki/%E7%BB%B4%E6%95%B0%E7%81%BE%E9%9A%BE>

<http://blog.csdn.net/zc02051126/article/details/49618633>

# K-Nearest Neighbor: Properties

- Advantages:

- Simple and intuitive, Training is very fast, easily implementable

- Particularly suitable for multi-classification problems

- With infinite training data and large enough K, K-NN approaches work well!

- Disadvantages:

- Sensitive to noisy features

- Store all the training data in memory even at test time

- Slow at query time:  $O(nd)$  computations for each test point

- In high dimensions, distance notions can be counter-intuitive!

- May perform badly in high dimensions (curse of dimensionality)

Also called:

- Memory/Instance-based learning

- Lazy learning

---


# Reducing Complexity

- Various exact and **approximate methods** for **reducing complexity**
- Computational complexity
  - use smart data structures, like **k-d trees**
  - ANN、BBF算法、LSH(局部敏感哈希)、Randomized K-d trees
  - 球树、M树、VP树、MVP树


# K-d Trees

- Invented in 1970s by Jon Bentley
- Name originally meant “3-d trees, 4-d trees, etc” where k was the # of dimensions
- Idea: Each level of the tree compares against 1 dimension.
- Tree used to store spatial data.
  - Nearest neighbor search.
  - Range queries.
  - Fast look-up!
- k-d trees are guaranteed  $\log_2 n$  depth where n is the number of points in the set.
  - Traditionally, k-d trees store points in d-dimensional space


# 3D k-d tree





## k-d tree construction(1)



## k-d tree construction(2)




## k-d tree construction(3)








## k-d tree construction(4)




k-d tree construction(5)




## k-d tree construction(6)




# k-d tree construction(7)




# k-d tree construction(8)




# k-d tree construction(9)



# k-d tree construction(10)





# k-d tree construction(11)









# k-d tree construction(12)





# k-d tree construction(13)





# k-d tree construction(14)





# k-d tree construction(15)




# k-d tree construction(16)



# k-d tree construction(17)



# k-d tree construction(18)



# Construction strategy

- The construction algorithm is similar as in 2-d
- At the root we split the set of points into two subsets of same size by a hyperplane vertical to  $x_1$ -axis
- At the children of the root, the partition is based on the second coordinate:  $x_2$ -coordinate
- At depth  $d$ , we start all over again by partitioning on the first coordinate
- The recursion stops until there is only one point left, which is stored as a leaf

Q1: Which dimension is used to split the set ?


axis with widest spread, biggest variance, alternating each one

Q2: The split point of the dimension

median, middle point of interval



# k-d Tree Splitting



sorted points in each dimension

	1	2	3	4	5	6	7	8	9
x	a	d	g	b	e	i	c	h	f
y	a	c	b	d	f	e	h	g	i

indicator for each set

	a	b	c	d	e	f	g	h	i
	0	0	1	0	0	1	0	1	1



scan sorted points in y dimension  
and add to correct set

y	a	b	d	e	g	c	f	h	i
---	---	---	---	---	---	---	---	---	---


Constructing the k-d tree can be done in  $O(d \log n)$  and  $O(dn)$  storage

# k-d Tree Nearest Neighbor Search


■ query point



# Why does k-d tree work?



$q(n.axis) - w \leq n.value$   
means the circle overlaps  
the left subtree.





$q(n.axis) + w > n.value$   
means the circle overlaps  
the right subtree.

## k-d Tree Nearest Neighbor Search

- Search recursively to find the point in the same cell as the query.
- On the return search each subtree where a closer point than the one you already know about might be found.



# k-d Tree NNS (1)

■ query point





## k-d Tree NNS (2)

■ query point





# k-d Tree NNS (3)

■ query point



# k-d Tree NNS (4)



■ query point







# k-d Tree NNS (5)

■ query point





# k-d Tree NNS (6)

■ query point





# k-d Tree NNS (7)

■ query point





# k-d Tree NNS (8)

■ query point





# k-d Tree NNS (9)

■ query point





# k-d Tree NNS (10)

■ query point





# k-d Tree NNS (11)

■ query point



# k-d Tree NNS (12)



■ query point







# k-d Tree NNS (13)

■ query point





# k-d Tree NNS (14)

■ query point





# k-d Tree NNS (15)

■ query point





# k-d Tree NNS (16)

■ query point





# k-d Tree NNS (17)

■ query point





# k-d Tree NNS (18)

■ query point





# k-d Tree NNS (19)

■ query point



# k-d Tree NNS (20)



■ query point





# k-d Tree NNS (21)

■ query point



# Nearest Neighbor Search

Main is `NNS(q,root,null,infinity)`

```
NNS(q: point, n: node, p: point, w: distance) : point {
  if n.left = null then {leaf case}
    if distance(q,n.point) < w then return n.point else return p;
  else
    if w = infinity then
      if q(n.axis) ≤ n.value then
        p := NNS(q,n.left,p,w);
        w := distance(p,q);
        if q(n.axis) + w > n.value then p := NNS(q, n.right, p, w);
      else
        p := NNS(q,n.right,p,w);
        w := distance(p,q);
        if q(n.axis) - w ≤ n.value then p := NNS(q, n.left, p, w);
    else //w is finite//
      if q(n.axis) - w ≤ n.value then
        p := NNS(q, n.left, p, w);
        w := distance(p,q);
        if q(n.axis) + w > n.value then p := NNS(q, n.right, p, w);
    return p
}
```

# complexity

- A data structure to support range queries in  $\mathbb{R}^d$
- Building a static  $k$ -d tree from  $n$  points takes  $O(dn \log n)$  time if sorting is used to compute the median at each level.
- Space complexity:  $O(dn)$
- Querying a balanced  $k$ -d tree takes  $O(n^{1-1/d} + k)$  time where  $k$  is the number of the reported points, and  $d$  is the dimension of the  $k$ -d tree.

K-D trees are **not suitable** for efficiently finding the nearest neighbor in **very high dimensional** spaces.

**Why?**

**Too many branches need to be backtracked! Close to linear time**

# Reducing Complexity

Find projection to a lower dimensional space so that the distances between samples are approximately the same

- PCA(Principal Component Analysis)
- Projection to a Random subspace
- ...

# Home work

- Hello World of Machine Learning!

- 1.实现K-近邻算法识别手写数字数据集。
- 2.改变K的值、修改为随机选取样本、改变训练样本数目，观察对算法错误率的影响。
- 3.体会“机器学习：数据驱动的科学”。



*Thanks !*